

Data Engineer Challenge

The goal of the challenge is to have a tool that is able to stream data from [kafka](#) and count unique things within this data. The simplest use case is that we want to calculate unique users per minute, day, week, month, year. [1] For a very first version, business wants us to provide just the unique users per minute. [1]

- The data consists of (Log)-Frames of JSON data that are streamed into/from apache kafka.
- Each frame has a timestamp property which is unixtime, the name of the property is `ts`.
- Each frame has a user id property called `uid`.
- You can assume that 99.9% of the frames arrive with a maximum latency of 5 seconds.
- you want to display the results as soon as possible
- the results should be forwarded to a new kafka topic (again as json.) choose a suitable structure. [1]
- for an advanced solution you should assume that you can *not* guarantee that events are always strictly ordered. [2]

Requirements:

- provide a Readme that shows how to build and run the code on Linux or OS X
- write a report: what did you do? what were the reasons you did it like that?
- measure at least one performance metric (e.g. frames per second)
- document your approach on how you decide **when** to output the data [1]
- document the estimated error in counting [3]
- it should be possible to ingest historical data. e.g. from the last 2 years. [1]

sample data

For a quick start you can use the sample data provided at <http://tx.tamedia.ch.s3.amazonaws.com/challenge/data/stream.jsonl.gz>:

it was generated using

```
./data-generator -c 1000000 -o stream.jsonl -r 1000 -n 1000000
```

```
cat stream.jsonl | jq .uid -r | gsort --parallel=4 | guniq | wc -l
```

yields 99993 unique ids.

you can just input this file to your kafka topic once by something like

```
gzcat stream.gz | kafka-console-producer --broker-list localhost:9092 --topic mytopic
```

You can also generate your own data by using the data-generator tool. it allows you to introduce random-ness if you have a more advanced solution and want to test it. it requires you to install a Dlang compiler and D's package manager. then `executedub build` to build the binary. binaries for os x and linux are in the bin directory.

suggested steps:

basic solution

1. install kafka
2. create a topic
3. use the kafka producer from kafka itself to send our test data to your topic
4. create a small app that reads this data from kafka and prints it to stdout
5. find a suitable data structure for counting [1] and implement a simple counting mechanism, output the results to stdout

advanced solution

6. benchmark
7. Output to a new Kafka Topic instead of stdout [1]
8. try to measure performance and optimize
9. write about how you could scale [1]
10. only now think about the edge cases, options and other things [3]

Evaluation Criteria & Expectations

- ability to break down *business* requirements into simple prototype code
- clean project setup and documentation
- research and use of suitable data structure for a specific use case. explain which and why. [3]
- ability to write performant code to handle streaming data. measure and document *how* fast your solution is.
- Understanding how to benchmark and analyze performance bottle necks. what tools did you use? [3]
- awareness of the mechanisms and costs of serialization. Explain (and ideally prove!) why json is an ideal format here or why not and then suggest a better solution.
- scalability: *explain* how you *would* scale your approach [2]

Bonus Questions / Challenges:

- how do you scale it to improve throughput. [1]
- you may want count things for different time frames but only do json parsing once. [1]
- explain how you would cope with failure if the app crashes mid day / mid year. [1]
- when creating e.g. per minute statistics, how do you handle frames that arrive late or frames with a random timestamp (e.g. hit by a bitflip), describe a strategy? [2]
- make it accept data also from std-in (instead of kafka) for rapid prototyping. (this might be helpful to have for testing anyways)
- measure also the performance impact / overhead of the json parser

Hints:

- tap into other peoples know-how and code.
- expected time is ~8 hours. If you are above that think about which parts to leave out and just document HOW you would do them.
- you should not use any big data framework, just your favourite fast programming language that has a Kafka driver. The most simplest version to achieve step 5 can be done in about 10 lines of code (plus the boilerplate to read from kafka).
- If you found the cheatcode for level 5 (well done!) you should try to achieve level 10 ;)
- check that your last commit compiles.
- be smart and impress us. It does not matter if you impress us with nice clean code or with a very clever hack to achieve the business goal in short time.